zoo FAQ

Gabor Grothendieck
GKX Associates Inc.

Achim Zeileis
Wirtschaftsuniversität Wien

Abstract

This is a collection of frequently asked questions (FAQ) about the **zoo** package together with their answers.

Keywords: irregular time series, daily data, weekly data, returns.

1. I know that duplicate times are not allowed but my data has them. What do I do?

zoo objects should not normally contain duplicate times. If you try to create such an object using **zoo** or `read.zoo` then warnings will be issued but the objects will be created. The user then has the opportunity to fix them up – typically by using `aggregate.zoo` or `duplicated`.

Merging is not well defined for duplicate series with duplicate times and rather than give an undesired or unexpected result, `merge.zoo` issues an error message if it encounters such illegal objects. Since `merge.zoo` is the workhorse behind many **zoo** functions, a significant portion of **zoo** will not accept duplicates among the times. Typically duplicates are eliminated by (1) averaging over them, (2) taking the last among each run of duplicates or (3) interpolating the duplicates and deleting ones on the end that cannot be interpolated. These three approaches are shown here using the `aggregate.zoo` function. Another way to do this is to use the `aggregate` argument of `read.zoo` which will aggregate the zoo object read in by `read.zoo` all in one step.

Note that in the example code below that `identity` is the identity function (i.e. it just returns its argument). It is an **R** core function:

A `"zoo"` series with duplicated indexes

```
> z <- suppressWarnings(zoo(1:8, c(1, 2, 2, 2, 3, 4, 5, 5)))
> z

1 2 2 2 3 4 5 5
1 2 3 4 5 6 7 8
```

Fix it up by averaging duplicates:

```
> aggregate(z, identity, mean)

  1   2   3   4   5
1.0 3.0 5.0 6.0 7.5
```

Or, fix it up by taking last in each set of duplicates:

```
> aggregate(z, identity, tail, 1)

1 2 3 4 5
1 4 5 6 8
```

Fix it up via interpolation of duplicate times

```
> time(z) <- na.approx(ifelse(duplicated(time(z)), NA, time(z)),
+     na.rm = FALSE)
```

If there is a run of equal times at end they wind up as NAs and we cannot have NA times.

```
> z[!is.na(time(z))]

    1       2 2.3333 2.6667       3       4       5
    1       2      3      4       5       6       7
```

The `read.zoo` command has an `aggregate` argument that supports arbitrary summarization. For example, in the following we take the last value among any duplicate times and sum the volumes among all duplicate times. We do this by reading the data twice, once for each aggregate function. In this example, the first three columns are junk that we wish to suppress which is why we specified `colClasses`; however, in most cases that argument would not be necessary.

```
> Lines <- "1|BHARTIARTL|EQ|18:15:05|600|1\n2|BHARTIARTL|EQ|18:15:05|600|99\n3|GLENMARK|EQ
> library("zoo")
> library("chron")
> tail1 <- function(x) tail(x, 1)
> cls <- c("NULL", "NULL", "NULL", "character", "numeric", "numeric")
> nms <- c("", "", "", "time", "value", "volume")
> z <- read.zoo(textConnection(Lines), aggregate = tail1, FUN = times,
+     sep = "|", colClasses = cls, col.names = nms)
> z2 <- read.zoo(textConnection(Lines), aggregate = sum, FUN = times,
+     sep = "|", colClasses = cls, col.names = nms)
> z$volume <- z2$volume
> z

         value volume
18:15:05  1100    217
18:15:06    80    201
```

2. When I try to specify a log axis to `plot.zoo` a warning is issued. What is wrong?

Arguments that are part of `...` are passed to the `panel` function and the default `panel` function, `lines`, does not accept `log`. Either ignore the warning, use `suppressWarnings` (see `?suppressWarnings`) or create your own panel function which excludes the `log`:

```
> z <- zoo(1:100)
> plot(z, log = "y", panel = function(..., log) lines(...))
```

3. How do I create right and a left vertical axes in plot.zoo?

The following shows an example of creating a plot containing a single panel and both left and right axes.

```
> set.seed(1)
> z.Date <- as.Date(paste(2003, 2, c(1, 3, 7, 9, 14), sep = "-"))
> z <- zoo(cbind(left = rnorm(5), right = rnorm(5, sd = 0.2)),
+     z.Date)
> plot(z[, 1], xlab = "Time", ylab = "")
> opar <- par(usr = c(par("usr")[1:2], range(z[, 2])))
> lines(z[, 2], lty = 2)
> axis(side = 4)
> legend("bottomright", lty = 1:2, legend = colnames(z), bty = "n")
> par(opar)
```


Figure 1: Left and right plot.zoo axes.

4. I have data frame with both numeric and factor columns. How do I convert that to a "zoo" object?

A "zoo" object may be (1) a numeric vector, (2) a numeric matrix or (3) a factor but may

not contain both a numeric vector and factor. You can do one of the following.
Use two `"zoo"` variables instead:

```
> DF <- data.frame(time = 1:4, x = 1:4, f = factor(letters[c(1,
+     1, 2, 2)]))
> zx <- zoo(DF$x, DF$time)
> zf <- zoo(DF$f, DF$time)
```

These could also be held in a `"data.frame"` again:

```
> DF2 <- data.frame(x = zx, f = zf)
```

Or convert the factor to numeric and create a single `"zoo"` series:

```
> z <- zoo(data.matrix(DF[-1]), DF$time)
```

5. Why does lag give slightly different results on a `"zoo"` and a `"zooreg"` series which are otherwise the same?

To be definite let us consider the following examples, noting how both `lag` and `diff` give a different answer with the same input except its class is `"zoo"` in one case and `"zooreg"` in another:

```
> z <- zoo(11:15, as.Date("2008-01-01") + c(-4, 1, 2, 3, 6))
> zr <- as.zooreg(z)
> lag(z)

2007-12-28 2008-01-02 2008-01-03 2008-01-04
        12         13         14         15

> lag(zr)

2007-12-27 2008-01-01 2008-01-02 2008-01-03 2008-01-06
        11         12         13         14         15

> diff(log(z))

2008-01-02 2008-01-03 2008-01-04 2008-01-07
0.08701138 0.08004271 0.07410797 0.06899287

> diff(log(zr))

2008-01-03 2008-01-04
0.08004271 0.07410797
```

`lag.zoo` and `lag.zooreg` work differently. For `"zoo"` objects the lagged version is obtained by moving values to the adjacent time point that exists in the series but for `"zooreg"` objects the time is lagged by `deltat`, the time between adjacent regular times.

A key implication is that `"zooreg"` can lag a point to a time point that did not previously exist in the series and, in particular, can lag a series outside of the original time range whereas that is not possible in a `"zoo"` series.

Note that `lag.zoo` has an `na.pad=` argument which in some cases may be what is being sought here.

The difference between `diff.zoo` and `diff.zooreg` stems from the fact that `diff(x)` is defined in terms of `lag` like this: `x-lag(x,-1)`.

6. How do I subtract the mean of each month from a `"zoo"` series?

Suppose we have a daily series. To subtract the mean of Jan 2007 from each day in that month, subtract the mean of Feb 2007 from each day in that month, etc. try this:

```
> set.seed(123)
> z <- zoo(rnorm(100), as.Date("2007-01-01") + seq(0, by = 10,
+     length = 100))
> z.demean1 <- z - ave(z, as.yearmon(time(z)))
```

This first generates some artificial data and then employs `ave` to compute monthly means.

To subtract the mean of all Januaries from each January, etc. try this:

```
> z.demean2 <- z - ave(z, format(time(z), "%m"))
```

7. How do I create a monthly series but still keep track of the dates?

Create a S3 subclass of `"yearmon"` called `"yearmon2"` that stores the dates as names on the time vector. It will be sufficient to create an `as.yearmon2` generic together with an `as.yearmon2.Date` methods as well as the inverse: `as.Date.yearmon2`.

```
> as.yearmon2 <- function(x, ...) UseMethod("as.yearmon2")
> as.yearmon2.Date <- function(x, ...) {
+     y <- as.yearmon(with(as.POSIXlt(x, tz = "GMT"), 1900 + year +
+         mon/12))
+     names(y) <- x
+     structure(y, class = c("yearmon2", class(y)))
+ }
```

`as.Date.yearmon2` is inverse of `as.yearmon2.Date`

```
> as.Date.yearmon2 <- function(x, frac = 0, ...) {
+     if (!is.null(names(x)))
+         return(as.Date(names(x)))
```

```
+       x <- unclass(x)
+       year <- floor(x + 0.001)
+       month <- floor(12 * (x - year) + 1 + 0.5 + 0.001)
+       dd.start <- as.Date(paste(year, month, 1, sep = "-"))
+       dd.end <- dd.start + 32 - as.numeric(format(dd.start + 32,
+           "%d"))
+       as.Date((1 - frac) * as.numeric(dd.start) + frac * as.numeric(dd.end),
+           origin = "1970-01-01")
+ }
```

This new class will act the same as `"yearmon"` stores and allows recovery of the dates using `as.Date` and `aggregate.zoo`.

```
> dd <- seq(as.Date("2000-01-01"), length = 5, by = 32)
> z <- zoo(1:5, as.yearmon2(dd))
> z
```

```
Jan 2000 Feb 2000 Mar 2000 Apr 2000 May 2000
       1        2        3        4        5
```

```
> aggregate(z, as.Date, identity)
```

```
2000-01-01 2000-02-02 2000-03-05 2000-04-06 2000-05-08
         1          2          3          4          5
```

8. How are axes added to a plot created using `plot.zoo`?

On single panel plots `axis` or `Axis` can be used just as with any classic graphics plot in R.

The following example adds custom axis for single panel plot. It labels months but uses the larger year for January. Months, quarters and years should have successively larger ticks.

```
> z <- zoo(0:500, as.Date(0:500))
> plot(z, xaxt = "n")
> tt <- time(z)
> m <- unique(as.Date(as.yearmon(tt)))
> jan <- format(m, "%m") == "01"
> mlab <- substr(months(m[!jan]), 1, 1)
> axis(side = 1, at = m[!jan], labels = mlab, tcl = -0.3, cex.axis = 0.7)
> axis(side = 1, at = m[jan], labels = format(m[jan], "%y"), tcl = -0.7)
> axis(side = 1, at = unique(as.Date(as.yearqtr(tt))), labels = FALSE)
> abline(v = m, col = grey(0.8), lty = 2)
```

A multivariate series can either be generated as (1) multiple single panel plots:

```
> z3 <- cbind(z1 = z, z2 = 2 * z, z3 = 3 * z)
> opar <- par(mfrow = c(2, 2))
```

```
> tt <- time(z)
> m <- unique(as.Date(as.yearmon(tt)))
> jan <- format(m, "%m") == "01"
> mlab <- substr(months(m[!jan]), 1, 1)
> for (i in 1:ncol(z3)) {
+     plot(z3[, i], xaxt = "n", ylab = colnames(z3)[i], ylim = range(z3))
+     axis(side = 1, at = m[!jan], labels = mlab, tcl = -0.3, cex.axis = 0.7)
+     axis(side = 1, at = m[jan], labels = format(m[jan], "%y"),
+         tcl = -0.7)
+     axis(side = 1, at = unique(as.Date(as.yearqtr(tt))), labels = FALSE)
+ }
> par(opar)
```

or (2) as a multipanel plot. In this case any custom axis must be placed in a panel function.

```
> plot(z3, screen = 1:3, xaxt = "n", nc = 2, ylim = range(z3),
+     panel = function(...) {
+         lines(...)
+         panel.number <- parent.frame()$panel.number
+         nser <- parent.frame()$nser
+         if (panel.number%%2 == 0 || panel.number == nser) {
+             tt <- list(...)[[1]]
+             m <- unique(as.Date(as.yearmon(tt)))
+             jan <- format(m, "%m") == "01"
+             mlab <- substr(months(m[!jan]), 1, 1)
+             axis(side = 1, at = m[!jan], labels = mlab, tcl = -0.3,
+                 cex.axis = 0.7)
+             axis(side = 1, at = m[jan], labels = format(m[jan],
+                 "%y"), tcl = -0.7)
+             axis(side = 1, at = unique(as.Date(as.yearqtr(tt))),
+                 labels = FALSE)
+         }
+     })
```

9. Why is nothing plotted except axes when I plot an object with many NAs?

Isolated points surrounded by NA values do not form lines:

```
> z <- zoo(c(1, NA, 2, NA, 3))
> plot(z)
```

So try one of the following:

Plot points rather than lines.

```
> plot(z, type = "p")
```

Omit `NA`s and plot that.

```
> plot(na.omit(z))
```

Fill in the `NA`s with interpolated values.

```
> plot(na.approx(z))
```

Plot points with lines superimposed.

```
> plot(z, type = "p")
> lines(na.omit(z))
```

Note that this is not specific to **zoo.** If we plot in R without **zoo** we get the same behavior.

10. Does zoo work with Rmetrics?

Yes. `timeDate` class objects from the **timeDate** package can be used directly as the index of a `zoo` series and `as.timeSeries.zoo` and `as.zoo.timeSeries` can convert back and forth between objects of class `zoo` and class `timeSeries` from the **timeSeries** package.

```
> library("timeDate")
> dts <- c("1989-09-28", "2001-01-15", "2004-08-30", "1990-02-09")
> tms <- c("23:12:55", "10:34:02", "08:30:00", "11:18:23")
> td <- timeDate(paste(dts, tms), format = "%Y-%m-%d %H:%M:%S")
> library("zoo")
> z <- zoo(1:4, td)
> zz <- merge(z, lag(z))
> plot(zz)
> library("timeSeries")
> zz

                    z lag.z.
1989-09-28 23:12:55 1      4
1990-02-09 11:18:23 4      2
2001-01-15 10:34:02 2      3
2004-08-30 08:30:00 3     NA

> as.timeSeries(zz)

GMT
                    z lag.z.
1989-09-28 23:12:55 1      4
1990-02-09 11:18:23 4      2
2001-01-15 10:34:02 2      3
2004-08-30 08:30:00 3     NA

> as.zoo(as.timeSeries(zz))
```

```
                  z lag.z.
1989-09-28 23:12:55 1      4
1990-02-09 11:18:23 4      2
2001-01-15 10:34:02 2      3
2004-08-30 08:30:00 3     NA
```

11. What other packages use zoo?

There are 53 other packages that depend on, suggest, can use with or are otherwise used with **zoo** that we have located:

Depends	
AER	Applied econometrics with R
BootPR	Bootstrap prediction intervals and bias-corrected fo
dyn	Time-series regression
dynlm	Dynamic linear regression
fda	Functional data analysis
FinTS	Companion to Tsay's "Analysis of financial time se
fractalrock	Generate fractal time series with non-normal retur bution
fUtilities	**Rmetrics** function utilities
fxregime	Exchange rate regime analysis
lmtest	Testing linear regression models
meboot Maximum Entropy Bootstrap for Time Series	
MFDF Modeling Functional Data in Finance	
party	Recursive partytioning toolbox
PerformanceAnalytics	Econometric tools for performance and risk analys
quantmod	Quantitative financial modelling framework
RBloomberg	R/**Bloomberg** interface
RcppTemplate	Rcpp R/C++ Object Mapping Library and Pack plate
sandwich	Robust covariance matrix estimators
sde	Simulation and Inference for Stochastic Different tions
StreamMetabollism	Calculating single station metabolism from diurna curves
strucchange	Testing, monitoring, and dating structural changes
tawny	Provides various portfolio optimization strategies random matrix theory and shrinkage estimators
tgram	Functions to compute and plot tracheidograms
tis	Regular time series package, previously part of fam
tripEstimation	Metropolis sampler and supporting functions for e animal movement from archival tags and satellite
tseries	Time series analysis and computational finance
TSfame	Time Series Database Interface extensions for fam
TShirtQuote	Time Series Database Interface extensions for get.h
VhayuR	R Interface to the Vhayu time series database
xts	Extensible time series
Suggests	
ChIPSim	Simulation of ChIP-seq experiments
futile	Utilities for strings and function arguments
gsubfn	Utilities for strings and function arguments
playwith	Interactive graphics: works with `xylot.zoo`
pscl	Political Science Computational Laboratory, Stan versity
tframePlus Time Frame coding kernel extensions	
TSdbi	Time series database interface
TSMySQL	Time series database interface extensions for **MyS**
TSPostgresSQL	Time series database interface extentions for **Postg**
TSodbc	Time series database interface extentions for **ODB**
TSSQLite	Time series database interface extentions for **SQLi**
wq	water quality monitoring
UsingR	A collection of datasets to accompany the textbo R for Introductory Statistics"
Zelig	Everyone's statistical software

12. Why does `ifelse` *not work as I expect?*

The ordinary R `ifelse` function only works with zoo objects if all three arguments are zoo objects with the same time index. `zoo` provides an `ifelse.zoo` function that should be used instead. The `.zoo` part must be written out since `ifelse` is not generic.

```
> z <- zoo(c(1, 5, 10, 15))
> ifelse(diff(z) > 4, -z, z)

  2   3   4
  1  -5 -10

> ifelse.zoo(diff(z) > 4, -z, z)

  1   2   3   4
 NA   5 -10 -15

> xm <- merge(z, dif = diff(z))
> with(xm, ifelse(dif > 4, -z, z))

  1   2   3   4
 NA   5 -10 -15

> ifelse(diff(z, na.pad = TRUE) > 4, -z, z)

  1   2   3   4
 NA   5 -10 -15
```

13. In a series which is regular except for a few missing times or for which we wish to align to a grid how is it filled or aligned?

```
> zym <- zoo(1:5, as.yearmon("2000-01-01") + c(0, 1, 2, 4, 5)/12)
> g <- seq(start(zym), end(zym), by = 1/12)
> na.locf(zym, xout = g)

Jan 2000 Feb 2000 Mar 2000 Apr 2000 May 2000 Jun 2000
       1        2        3        3        4        5
```

A variation of this is where the grid is of a different date/time class than the original series. In that case use the x argument. In the example that follows the series `z` is of `"Date"` class whereas the grid is of `"yearmon"` class:

```
> z <- zoo(1:3, as.Date(c("2000-01-15", "2000-03-3", "2000-04-29")))
> g <- seq(as.yearmon(start(z)), as.yearmon(end(z)), by = 1/12)
> na.locf(z, x = as.yearmon, xout = g)
```

```
Jan 2000 Feb 2000 Mar 2000 Apr 2000
       1        1        2        3
```

Here is a `chron` example where we wish to create a 10 minute grid:

```
> Lines <- "Time,Value\n2009-10-09 5:00:00,210\n2009-10-09 5:05:00,207\n2009-10-09 5:17:00
> library(chron)
> z <- read.zoo(textConnection(Lines), FUN = as.chron, sep = ",",
+     header = TRUE)
> g <- seq(start(z), end(z), by = times("00:10:00"))
> na.locf(z, xout = g)

(10/09/09 05:00:00) (10/09/09 05:10:00) (10/09/09 05:20:00) (10/09/09 05:30:00)
                210                 207                 250                 193
(10/09/09 05:40:00) (10/09/09 05:50:00) (10/09/09 06:00:00)
                193                 205                 185
```

Affiliation:

Gabor Grothendieck
GKX Associates Inc.
E-mail: ggrothendieck@gmail.com

Achim Zeileis
Wirtschaftsuniversität Wien
E-mail: Achim.Zeileis@R-project.org