zoo FAQ

Gabor Grothendieck
GKX Associates Inc.

Achim Zeileis
Wirtschaftsuniversität Wien

Abstract

This is a collection of frequently asked questions (FAQ) about the **zoo** package together with their answers.

Keywords: irregular time series, daily data, weekly data, returns.

1. I know that duplicate times are not allowed but my data has them. What do I do?

zoo objects should not normally contain duplicate times. If you try to create such an object using **zoo** or `read.zoo` then warnings will be issued but the objects will be created. The user then has the opportunity to fix them up – typically by using `aggregate.zoo` or `duplicated`.

Merging is not well defined for duplicate series with duplicate times and rather than give an undesired or unexpected result, `merge.zoo` issues an error message if it encounters such illegal objects. Since `merge.zoo` is the workhorse behind many **zoo** functions, a significant portion of **zoo** will not accept duplicates among the times. Typically duplicates are eliminated by (1) averaging over them, (2) taking the last among each run of duplicates or (3) interpolating the duplicates and deleting ones on the end that cannot be interpolated. These three approaches are shown here using the `aggregate.zoo` function. Another way to do this is to use the `aggregate` argument of `read.zoo` which will aggregate the zoo object read in by `read.zoo` all in one step.

Note that in the example code below that `force` is the identity function (i.e. it just returns its argument). It is an **R** core function:

A "zoo" series with duplicated indexes

```
> z <- suppressWarnings(zoo(1:8, c(1, 2, 2, 2, 3, 4, 5, 5)))
> z


1 2 2 2 3 4 5 5
1 2 3 4 5 6 7 8
```

Fix it up by averaging duplicates:

```
> aggregate(z, force, mean)


  1   2   3   4   5
1.0 3.0 5.0 6.0 7.5
```

Or, fix it up by taking last in each set of duplicates:

```
> aggregate(z, force, tail, 1)

1 2 3 4 5
1 4 5 6 8
```

Fix it up via interpolation of duplicate times

```
> time(z) <- na.approx(ifelse(duplicated(time(z)), NA, time(z)),
+     na.rm = FALSE)
```

If there is a run of equal times at end they wind up as NAs and we cannot have NA times.

```
> z[!is.na(time(z))]

    1      2 2.3333 2.6667      3      4      5
    1      2      3      4      5      6      7
```

The `read.zoo` command has an `aggregate` argument that supports arbitrary summarization. For example, in the following we take the last value among any duplicate times and sum the volumes among all duplicate times. We do this by reading the data twice, once for each aggregate function. In this example, the first three columns are junk that we wish to suppress which is why we specified `colClasses`; however, in most cases that argument would not be necessary.

```
> Lines <- "1|BHARTIARTL|EQ|18:15:05|600|1\n2|BHARTIARTL|EQ|18:15:05|600|99\n3|GLENMARK|EQ
> library("zoo")
> library("chron")
> tail1 <- function(x) tail(x, 1)
> cls <- c("NULL", "NULL", "NULL", "character", "numeric", "numeric")
> nms <- c("", "", "", "time", "value", "volume")
> z <- read.zoo(textConnection(Lines), aggregate = tail1, FUN = times,
+     sep = "|", colClasses = cls, col.names = nms)
> z2 <- read.zoo(textConnection(Lines), aggregate = sum, FUN = times,
+     sep = "|", colClasses = cls, col.names = nms)
> z$volume <- z2$volume
> z

         value volume
18:15:05  1100    217
18:15:06    80    201
```

2. *When I try to specify a log axis to* `plot.zoo` *a warning is issued. What is wrong?*

Arguments that are part of `...` are passed to the `panel` function and the default `panel` function, `lines`, does not accept `log`. Either ignore the warning, use `suppressWarnings` (see `?suppressWarnings`) or create your own panel function which excludes the `log`:

```
> z <- zoo(1:100)
> plot(z, log = "y", panel = function(..., log) lines(...))
```

3. How do I create right and a left vertical axes in plot.zoo?

The following shows an example of creating a plot containing a single panel and both left and right axes.

```
> set.seed(1)
> z.Date <- as.Date(paste(2003, 2, c(1, 3, 7, 9, 14), sep = "-"))
> z <- zoo(cbind(left = rnorm(5), right = rnorm(5, sd = 0.2)),
+     z.Date)
> plot(z[, 1], xlab = "Time", ylab = "")
> opar <- par(usr = c(par("usr")[1:2], range(z[, 2])))
> lines(z[, 2], lty = 2)
> axis(side = 4)
> legend("bottomright", lty = 1:2, legend = colnames(z), bty = "n")
> par(opar)
```


Figure 1: Left and right plot.zoo axes.

4. I have data frame with both numeric and factor columns. How do I convert that to a "zoo" object?

A "zoo" object may be (1) a numeric vector, (2) a numeric matrix or (3) a factor but may

not contain both a numeric vector and factor. You can do one of the following.
Use two "zoo" variables instead:

```
> DF <- data.frame(time = 1:4, x = 1:4, f = factor(letters[c(1,
+     1, 2, 2)]))
> zx <- zoo(DF$x, DF$time)
> zf <- zoo(DF$f, DF$time)
```

These could also be held in a "data.frame" again:

```
> DF2 <- data.frame(x = zx, f = zf)
```

Or convert the factor to numeric and create a single "zoo" series:

```
> z <- zoo(data.matrix(DF[-1]), DF$time)
```

5. Why does lag give slightly different results on a "zoo" and a "zooreg" series which are otherwise the same?

To be definite let us consider the following examples, noting how both `lag` and `diff` give a different answer with the same input except its class is "zoo" in one case and "zooreg" in another:

```
> z <- zoo(11:15, as.Date("2008-01-01") + c(-4, 1, 2, 3, 6))
> zr <- as.zooreg(z)
> lag(z)

2007-12-28 2008-01-02 2008-01-03 2008-01-04
        12         13         14         15

> lag(zr)

2007-12-27 2008-01-01 2008-01-02 2008-01-03 2008-01-06
        11         12         13         14         15

> diff(log(z))

2008-01-02 2008-01-03 2008-01-04 2008-01-07
0.08701138 0.08004271 0.07410797 0.06899287

> diff(log(zr))

2008-01-03 2008-01-04
0.08004271 0.07410797
```

`lag.zoo` and `lag.zooreg` work differently. For `"zoo"` objects the lagged version is obtained by moving values to the adjacent time point that exists in the series but for `"zooreg"` objects the time is lagged by `deltat`, the time between adjacent regular times.

A key implication is that `"zooreg"` can lag a point to a time point that did not previously exist in the series and, in particular, can lag a series outside of the original time range whereas that is not possible in a `"zoo"` series.

Note that `lag.zoo` has an `na.pad=` argument which in some cases may be what is being sought here.

The difference between `diff.zoo` and `diff.zooreg` stems from the fact that `diff(x)` is defined in terms of `lag` like this: `x-lag(x,-1)`.

6. How do I subtract the mean of each month from a `"zoo"` series?

Suppose we have a daily series. To subtract the mean of Jan 2007 from each day in that month, subtract the mean of Feb 2007 from each day in that month, etc. try this:

```
> set.seed(123)
> z <- zoo(rnorm(100), as.Date("2007-01-01") + seq(0, by = 10,
+     length = 100))
> z.demean1 <- z - ave(z, as.yearmon(time(z)))
```

This first generates some artificial data and then employs `ave` to compute monthly means.

To subtract the mean of all Januaries from each January, etc. try this:

```
> z.demean2 <- z - ave(z, format(time(z), "%m"))
```

7. How do I create a monthly series but still keep track of the dates?

Create a S3 subclass of `"yearmon"` called `"yearmon2"` that stores the dates as names on the time vector. It will be sufficient to create an `as.yearmon2` generic together with an `as.yearmon2.Date` methods as well as the inverse: `as.Date.yearmon2`.

```
> as.yearmon2 <- function(x, ...) UseMethod("as.yearmon2")
> as.yearmon2.Date <- function(x, ...) {
+     y <- as.yearmon(with(as.POSIXlt(x, tz = "GMT"), 1900 + year +
+         mon/12))
+     names(y) <- x
+     structure(y, class = c("yearmon2", class(y)))
+ }
```

`as.Date.yearmon2` is inverse of `as.yearmon2.Date`

```
> as.Date.yearmon2 <- function(x, frac = 0, ...) {
+     if (!is.null(names(x)))
+         return(as.Date(names(x)))
```

```
+       x <- unclass(x)
+       year <- floor(x + 0.001)
+       month <- floor(12 * (x - year) + 1 + 0.5 + 0.001)
+       dd.start <- as.Date(paste(year, month, 1, sep = "-"))
+       dd.end <- dd.start + 32 - as.numeric(format(dd.start + 32,
+           "%d"))
+       as.Date((1 - frac) * as.numeric(dd.start) + frac * as.numeric(dd.end),
+           origin = "1970-01-01")
+ }
```

This new class will act the same as `"yearmon"` stores and allows recovery of the dates using `as.Date` and `aggregate.zoo`.

```
> dd <- seq(as.Date("2000-01-01"), length = 5, by = 32)
> z <- zoo(1:5, as.yearmon2(dd))
> z
```

```
Jan 2000 Feb 2000 Mar 2000 Apr 2000 May 2000
       1        2        3        4        5
```

```
> aggregate(z, as.Date, force)
```

```
2000-01-01 2000-02-02 2000-03-05 2000-04-06 2000-05-08
         1          2          3          4          5
```

8. How are axes added to a plot created using `plot.zoo`?

On single panel plots `axis` or `Axis` can be used just as with any classic graphics plot in R.

The following example adds custom axis for single panel plot. It labels months but uses the larger year for January. Months, quarters and years should have successively larger ticks.

```
> z <- zoo(0:500, as.Date(0:500))
> plot(z, xaxt = "n")
> tt <- time(z)
> m <- unique(as.Date(as.yearmon(tt)))
> jan <- format(m, "%m") == "01"
> mlab <- substr(months(m[!jan]), 1, 1)
> axis(side = 1, at = m[!jan], labels = mlab, tcl = -0.3, cex.axis = 0.7)
> axis(side = 1, at = m[jan], labels = format(m[jan], "%y"), tcl = -0.7)
> axis(side = 1, at = unique(as.Date(as.yearqtr(tt))), labels = FALSE)
> abline(v = m, col = grey(0.8), lty = 2)
```

A multivariate series can either be generated as (1) multiple single panel plots:

```
> z3 <- cbind(z1 = z, z2 = 2 * z, z3 = 3 * z)
> opar <- par(mfrow = c(2, 2))
```

```
> tt <- time(z)
> m <- unique(as.Date(as.yearmon(tt)))
> jan <- format(m, "%m") == "01"
> mlab <- substr(months(m[!jan]), 1, 1)
> for (i in 1:ncol(z3)) {
+     plot(z3[, i], xaxt = "n", ylab = colnames(z3)[i], ylim = range(z3))
+     axis(side = 1, at = m[!jan], labels = mlab, tcl = -0.3, cex.axis = 0.7)
+     axis(side = 1, at = m[jan], labels = format(m[jan], "%y"),
+         tcl = -0.7)
+     axis(side = 1, at = unique(as.Date(as.yearqtr(tt))), labels = FALSE)
+ }
> par(opar)
```

or (2) as a multipanel plot. In this case any custom axis must be placed in a panel function.

```
> plot(z3, screen = 1:3, xaxt = "n", nc = 2, ylim = range(z3),
+     panel = function(...) {
+         lines(...)
+         panel.number <- parent.frame()$panel.number
+         nser <- parent.frame()$nser
+         if (panel.number%%2 == 0 || panel.number == nser) {
+             tt <- list(...)[[1]]
+             m <- unique(as.Date(as.yearmon(tt)))
+             jan <- format(m, "%m") == "01"
+             mlab <- substr(months(m[!jan]), 1, 1)
+             axis(side = 1, at = m[!jan], labels = mlab, tcl = -0.3,
+                 cex.axis = 0.7)
+             axis(side = 1, at = m[jan], labels = format(m[jan],
+                 "%y"), tcl = -0.7)
+             axis(side = 1, at = unique(as.Date(as.yearqtr(tt))),
+                 labels = FALSE)
+         }
+     })
```

9. Why is nothing plotted except axes when I plot an object with many NAs?

Isolated points surrounded by `NA` values do not form lines:

```
> z <- zoo(c(1, NA, 2, NA, 3))
> plot(z)
```

So try one of the following:

Plot points rather than lines.

```
> plot(z, type = "p")
```

Omit `NA`s and plot that.

```
> plot(na.omit(z))
```

Fill in the `NA`s with interpolated values.

```
> plot(na.approx(z))
```

Plot points with lines superimposed.

```
> plot(z, type = "p")
> lines(na.omit(z))
```

Note that this is not specific to **zoo.** If we plot in R without **zoo** we get the same behavior.

10. Does zoo work with Rmetrics?

Yes. `timeDate` class objects from the **timeDate** package can be used directly as the index of a `zoo` series and `as.timeSeries.zoo` and `as.zoo.timeSeries` can convert back and forth between objects of class `zoo` and class `timeSeries` from the **timeSeries** package.

```
> library("timeDate")
> dts <- c("1989-09-28", "2001-01-15", "2004-08-30", "1990-02-09")
> tms <- c("23:12:55", "10:34:02", "08:30:00", "11:18:23")
> td <- timeDate(paste(dts, tms), format = "%Y-%m-%d %H:%M:%S")
> library("zoo")
> z <- zoo(1:4, td)
> zz <- merge(z, lag(z))
> plot(zz)
> library("timeSeries")
> zz

                    z lag(z)
1989-09-28 23:12:55 1      4
1990-02-09 11:18:23 4      2
2001-01-15 10:34:02 2      3
2004-08-30 08:30:00 3     NA

> as.timeSeries(zz)

GMT
                    z lag(z)
1989-09-28 23:12:55 1      4
1990-02-09 11:18:23 4      2
2001-01-15 10:34:02 2      3
2004-08-30 08:30:00 3     NA

> as.zoo(as.timeSeries(zz))
```

```
                   z lag(z)
1989-09-28 23:12:55 1      4
1990-02-09 11:18:23 4      2
2001-01-15 10:34:02 2      3
2004-08-30 08:30:00 3     NA
```

11. What other packages use zoo?

Depends	
AER	Applied econometrics with R
BootPR	Bootstrap prediction intervals and bias-corrected forecasting
dyn	Time-series regression
dynlm	Dynamic linear regression
fda	Functional data analysis
FinTS	Companion to Tsay's "Analysis of financial time series"
fUtilities	**Rmetrics** function utilities
fxregime	Exchange rate regime analysis
lmtest	Testing linear regression models
party	Recursive partytioning toolbox
PerformanceAnalytics	Econometric tools for performance and risk analysis
quantmod	Quantitative financial modelling framework
RBloomberg	R/**Bloomberg** interface
sandwich	Robust covariance matrix estimators
strucchange	Testing, monitoring, and dating structural changes
tis	Regular time series package, previously part of fame package
tripEstimation	Metropolis sampler and supporting functions for estimating animal movement from archival tags and satellite fixes
tseries	Time series analysis and computational finance
VhayuR	R Interface to the Vhayu time series database
xts	Extensible time series
Suggests	
gsubfn	Utilities for strings and function arguments
pscl	Political Science Computational Laboratory, Stanford University
TSSQLite	Time series database interface extentions for **SQLite**
TSdbi	Time series database interface
Zelig	Everyone's statistical software
Uses or Used with	
timeDate	**Rmetrics** date and time functions: `timeDate` usable with `zoo`
grid	Graphics infrastructure: use with `xyplot.zoo`
its	Irregular time series: `as.its.zoo`, `as.zoo.its`
lattice	**grid**-based graphics: use with `xyplot.zoo`
playwith	Interactive graphics: works with `xylot.zoo`
timeSeries	**Rmetrics** time series functions: `as.timeSeries.zoo`, `as.zoo.timeSeries`
YaleToolkit	Data exploration tools from Yale University: accepts `"zoo"` input

Why does `ifelse` not work as I expect?

The ordinary R `ifelse` function only works with zoo objects if all three arguments are zoo objects with the same time index. **zoo** provides an `ifelse.zoo` function that should be used

instead. The `.zoo` part must be written out since `ifelse` is not generic.

```
> z <- zoo(c(1, 5, 10, 15))
> ifelse(diff(z) > 4, -z, z)

  2   3   4
  1  -5 -10

> ifelse.zoo(diff(z) > 4, -z, z)

  1   2   3   4
 NA   5 -10 -15

> xm <- merge(z, dif = diff(z))
> with(xm, ifelse(dif > 4, -z, z))

  1   2   3   4
 NA   5 -10 -15

> ifelse(diff(z, na.pad = TRUE) > 4, -z, z)

  1   2   3   4
 NA   5 -10 -15
```

Affiliation:

Gabor Grothendieck
GKX Associates Inc.
E-mail: ggrothendieck@gmail.com

Achim Zeileis
Wirtschaftsuniversität Wien
E-mail: Achim.Zeileis@R-project.org